

CoalSafe-MMTS: Multimodal Time-Series AI System

Sub-Surface Coal Spontaneous Combustion — Complete Technical Report v2.0

Student: Manas PurTTY | **Model:** CoalSafe-MMTS (653,329 parameters) | **GPU:** NVIDIA GeForce GTX 1650 (CUDA)

GitHub: <https://github.com/ManasPurTTY/CoalSafe-Sim-v2>

1. Executive Summary

Coal spontaneous combustion begins **deep underground** through low-temperature oxidation. Conventional monitoring is **reactive** — alarms trigger only after surface heat is visible, by which time the internal core has already reached critical runaway ($>100^{\circ}\text{C}$).

CoalSafe-MMTS solves this with 3 innovations:

- **Multimodal Fusion:** Combines 1-hour gas/temperature sensor sequences (Bi-GRU + Attention) with 50×50 infrared camera images (2D-CNN) for comprehensive situational awareness.
- **+1-Hour Future Forecasting:** Predicts future risk score and internal core temperature 60 minutes ahead — enabling pre-emptive cooling *before* fire develops.
- **Cross-Modal Conflict Gating:** Validates cross-sensor consensus before triggering any automated mitigation, suppressing false alarms from sensor glitches or solar camera glare.

2. Neural Network Architecture (CoalSafe-MMTS)

Block 1 — Temporal Sensor Branch (Bi-GRU + Temporal Attention)

Input: 6 time steps $\times$ 10 sensor channels = 60-minute sliding history window.

Architecture: 2-layer Bidirectional GRU (hidden=64) $\rightarrow$ Temporal Self-Attention $\rightarrow$ 128-d vector h_{tab} .

Block 2 — Spatial Thermal Image Branch (2D-CNN)

Input: 50 $\times$ 50 $\times$ 3 infrared thermal heatmap image.

Architecture: 3-stage Conv2D (32 $\rightarrow$ 64 $\rightarrow$ 128 filters, LeakyReLU) $\rightarrow$ Adaptive Pool $\rightarrow$ 128-d vector h_{img} .

Block 3 — Gated Bilinear Cross-Modal Fusion

Adaptive gate $g = \sigma(W_g [h_{\text{tab}}; h_{\text{img}}]) \rightarrow$ fused 256-d representation. Automatically weights sensor vs. camera reliability per moment.

Block 4 — Multi-Horizon Prediction Heads (NEW)

- `fused_risk_now` — Present-moment unified risk score (%)
- `risk_forecast_30m` — Risk score in +30 minutes
- `risk_forecast_60m` — **Risk score in +60 minutes (+1-Hour Forecast)**
- `temp_core_forecast_60m` — **Internal core temperature in +60 minutes (°C)**
- `class_logits_60m` — 5-class future category: NORMAL / LOW / MEDIUM / HIGH / CRITICAL




Figure 1: Multi-Task Training Loss Convergence over 40 epochs on NVIDIA GTX 1650 CUDA GPU. Best +1h Forecast MAE: 4.03% (achieved at Epoch 17).

3. +1-Hour Future Forecasting: Test Results

Model was tested on held-out Scenario S07 (never seen during training). The +1h forecasting head correctly tracks the future risk trajectory with an R^2 of **0.6926**.




Figure 2: Walk-forward test on Scenario S07 — AI +1h Predicted Risk (orange) vs Ground Truth Future Risk (cyan) and internal core temperature (purple).

4. Predictive Early Mitigation vs. Reactive Mitigation

Three operational policies evaluated over 24 hours (1,440 minutes):

Policy	Action Time	Peak Surface Temp	Peak Core Temp	Outcome
S06: AI Proactive Early (Our System)	t=360 min (6h)	31.2°C	38.1°C	★ 100% Runaway Averted
S07: Reactive Late (Traditional)	t=840 min (14h)	98.8°C	92.2°C	Severe Core Damage (+8h delay)
S05: Unmitigated Runaway	None	217.8°C	200.0°C	Total Combustion Loss




Figure 3: Internal core temperature & cumulative thermal damage comparing Proactive Early (green), Late Reactive (orange), and Unmitigated (red) policies over 24 hours.

5. Cross-Modal Conflict Gating & False Alarm Protection

Disagreement $|\Delta| = |y_{\text{sensor}} - y_{\text{camera}}|$. If $|\Delta| > 25\%$, the system flags **FALSE (Conflict)** and blocks automated spray.

Case	Condition	Sensor Risk	Camera Risk	$ \Delta $	Status	Action
Case 1: Normal	Low CO, cold camera	4.7%	5.7%	3.3%	TRUE (Consensus)	Routine Monitoring
Case 2: Developing Runaway	Rising CO + hot hotspot	70.4%	67.9%	0.8%	TRUE (Consensus)	★ Early Spray Triggered
Case 3: Sensor Glitch	CO spike (250ppm), cold camera	60.7%	8.6%	52.1%	FALSE (Conflict)	Spray Suppressed
Case 4: Camera Glare	Normal gas, sunlight glare	5.3%	66.5%	61.2%	FALSE (Conflict)	Spray Suppressed




Figure 4: Cross-Modal Conflict Gating Benchmark — 100% false-alarm suppression on Cases 3 & 4, with correct consensus validation on Cases 1 & 2.

6. Forecast Validation & Walk-Forward Backtesting

To prove the trustworthiness of the +1h forecast, we used **Walk-Forward Backtesting** on two completely held-out scenarios (S07, S08 — never seen during training), using the 24-hour simulation data as ground truth. At every time step t , the model predicted risk at $t+60\text{ min}$ and the prediction was compared against the actual simulation value.

Standard Industry Evaluation Metrics

Metric	What It Measures	S07 Result	S08 Result	Benchmark
MAE	Average absolute prediction error (%)	8.54%	5.69%	Good (<10%)
RMSE	Error (large mistakes penalized more)	11.20%	7.60%	Acceptable (<15%)
MAPE	% error relative to actual value	32.4%	47.8%	High (near-zero actuals inflate this)
R² Score	Variance explained (0→1)	0.6873	0.3312	S07: Good (>0.60)
Directional Accuracy	UP/DOWN trend prediction	62.1%	62.1%	Above random (50%)
Hit Rate	Real dangers correctly detected	60.0%	100.0%	Target: >70%
Miss Rate ★	Real dangers MISSED (critical!)	40.0%	0.0%	Target: 0%
False Alarm Rate	Safe moments falsely alarmed	17.8%	0.0%	Target: <10%
F1-Score	Harmonic mean of Precision & Recall	18.75	100.0	S08: Perfect
95% Band Coverage	Actual inside predicted range	94.3%	97.1%	Target: ~95% ✓
Horizon Degradation	Accuracy loss: +30min → +60min	+1.8% MAE	+5.2% MAE	Minimal degradation ✓




Figure 5: Walk-Forward Backtesting on S07 — Predicted vs Actual +1h Risk, per-step error bars, directional accuracy, and safety confusion matrix.




Figure 6: Walk-Forward Backtesting on S08 (Sensor Fault Scenario) — 100% Hit Rate, 0% Miss Rate, 0% False Alarms. Perfect classification under noisy conditions.



Figure 7: Multi-Scenario Evaluation Metrics Bar Chart comparing S07 and S08 across MAE, RMSE, R^2 , Hit Rate, Miss Rate, and Directional Accuracy.

7. Codebase Structure

File	Role	Key Output
code/ml/dataset.py	Multimodal time-series data loader with sliding 60-min window	Multi-horizon targets: y_now, y_30m, y_60m, T_core_60m
code/ml/model.py	Neural network: Bi-GRU + CNN + Gated Fusion + Multi-Head outputs	653,329 trainable parameters
code/ml/train.py	GPU training with multi-task loss (risk + forecast + temperature)	Best +1h MAE: 4.03% at Epoch 17
code/ml/evaluate.py	4-case conflict gating benchmark + held-out test metrics	evaluation_results.json
code/ml/predictive_controller.py	Closed-loop simulation: S06 vs S07 vs S05 thermal comparison	Lead-time advantage analysis
code/ml/forecast_validator.py	Walk-forward backtesting with 9 standard field metrics	forecast_validation_report.json + 3 validation graphs
code/ml/plot_ml_results.py	Generates all publication-quality result figures	8 graphs in graphs/ml_results/

8. Professor / Viva Examination Q&A

Q1: Why is +1h future forecasting better than detecting current risk?

Coal has high thermal inertia. By the time surface alarms trigger, the sub-surface core is already in runaway. Forecasting +1h ahead lets us cool the pile at 38°C instead of reacting at 92°C. Our system averts 100% of thermal runaway in proactive mode (S06).

Q2: How do you validate the forecast without real-world data?

We use Walk-Forward Backtesting on held-out simulation scenarios (S07, S08). At every time step t , we predict risk at $t+60m$ and compare against the actual simulation value. We evaluate using 9 standard metrics: MAE, RMSE, MAPE, R^2 , Directional Accuracy, Hit Rate, Miss Rate, False Alarm Rate, and 95% Confidence Band Coverage.

Q3: How does the model know internal core temperature without a physical probe?

The Bi-GRU with Temporal Attention learns the thermodynamic transfer function from surface gas accumulation rates (ΔCO , ΔCO_2 , heating rate) over the 60-minute sliding window and maps this to sub-surface core temperature predictions via the `temp_core_forecast_60m` output head.

Q4: How does the Conflict Gate prevent false spray activations?

The model computes separate risk estimates from sensors (y_{tab}) and camera (y_{img}). If $|y_{\text{tab}} - y_{\text{img}}| > 25\%$, the gate flags FALSE (Conflict) and blocks spray activation — catching CO electrical spikes (Case 3) and solar lens glare (Case 4) with 100% accuracy in testing.